

MarketPilot AI — Product Requirements Document (PRD)

Version: 1.0 **Prepared for:** Development Team **Product Type:** AI-Powered Marketing Consulting Platform **Document Owner:** Akshat Singh Bisht

1. Product Overview

MarketPilot AI is an AI-driven marketing consulting platform that simulates a full-service consulting firm. Instead of answering isolated questions, it runs a structured, multi-stage consulting engagement for any business — collecting inputs, researching the market/competitors/customers, and producing a complete strategic report with an execution roadmap.

Core differentiator: Most AI marketing tools generate generic content (a blog outline, an ad headline). MarketPilot AI generates a *structured consulting engagement* — discovery → research → analysis → strategy → execution plan — the way a real agency would, but compressed into minutes instead of weeks.

2. Problem Statement

Problem	Who feels it	Current alternative
Small/mid businesses can't afford a real marketing agency (₹50K-₹5L/month)	SMB owners, solopreneurs	DIY guesswork, generic YouTube advice
Freelance marketers/consultants spend hours on research before they can advise a client	Marketing consultants, freelancers	Manual research across 10+ tools
Generic AI chatbots give shallow, unstructured marketing advice	Anyone using ChatGPT for marketing help	One-off prompts, no continuity or structure
No single tool takes a business from "zero context" to a "ready-to-execute plan"	Startups, agencies	Piecing together templates, consultants, spreadsheets

3. Target Users

- Primary:** Small business owners / solopreneurs / D2C founders needing a marketing strategy without hiring an agency.

2. **Secondary:** Freelance marketing consultants who use the tool to accelerate client work (this is likely Akshat's own use case as an independent digital marketing researcher).
 3. **Tertiary:** Marketing agencies wanting an internal research/strategy co-pilot for junior staff.
-

4. Product Vision & Goals

Vision: Make world-class marketing strategy accessible to any business, regardless of budget, by replicating the output of a 25+ person consulting team through AI.

V1 Goals:

- Deliver a complete, structured strategy report from a guided intake conversation.
 - Never fabricate data — flag assumptions and ask for missing information.
 - Output must be immediately actionable by a non-marketer.
 - Reports must be exportable (PDF/Doc) and shareable.
-

5. Core Concept — Simulated Consulting Firm

The product does not behave as a single chatbot persona. It behaves as a firm with role-specialized "experts" whose outputs are merged into one report. This matters for prompt architecture and UI — each section of the report should visibly map to a "department" so the user trusts the depth (e.g., "SEO Consultant," "Consumer Psychology Expert").

Roster (for reference — drives section ownership in the report, not literal multi-agent chat): CEO Strategy Consultant, Business Consultant, Startup Advisor, Brand Strategist, Market Research Analyst, Competitive Intelligence Expert, Consumer Psychology Expert, Customer Research Specialist, Marketing Consultant, Digital Marketing Expert, Performance Marketing Specialist, Google/Facebook/LinkedIn Ads Experts, SEO Consultant (+ Local, Technical), Content Marketing Strategist, Copywriter, Social Media Strategist, Email Marketing Specialist, CRM Consultant, Sales Funnel Expert, CRO Specialist, Pricing Strategist, Product Marketing Manager, Data Analyst, BI Analyst, Growth Hacker, Customer Success Consultant.

6. End-to-End Workflow (Product Logic — Must Not Be Skipped)

Step	Name	Description	Blocking?
1	Understand the Business	Parse whatever the user has already provided	No

Step	Name	Description	Blocking?
2	Collect Missing Information	Ask targeted follow-up questions (see §7)	Yes — cannot proceed to research without core fields
3	Market Research	Industry size, trends, demand, gaps	No
4	Competitor Research	Direct/indirect/local/global competitors	No
5	Customer Research	Personas, psychographics, journey	No
6	Identify Opportunities	Cross-reference gaps from steps 3-5	No
7	Identify Weaknesses	SWOT/PESTLE/Porter's inputs	No
8	Create Strategy	Brand, marketing, sales, pricing, GTM	No
9	Execution Roadmap	30/60/90-day + 6/12-month plans	No
10	Estimate Outcomes	Expected ROI, KPI targets, ranges with caveats	No

Product rule: Step 2 is a hard gate. The UI should not allow report generation until required fields are filled or explicitly marked "unknown — proceed with assumption flagged."

7. Input Data Schema (Business Discovery Form)

This is the intake form / conversational data model. Recommend building this as a structured form (not free chat) for reliability, with an optional "chat mode" that maps answers back into these fields.

Required fields (block generation if missing):

- Business name
- Industry / category
- Products/services offered
- Business model (B2B/B2C/D2C/marketplace/SaaS/etc.)
- Business stage (idea / early / established / scaling)
- Target country & cities
- Target audience (basic description)
- Current monthly budget for marketing

- Primary goal for this engagement

Optional but high-value fields:

- Website URL (used for auto-scraping basic signals — traffic estimate, tech stack, SEO snapshot)
- Revenue model, current revenue range, AOV, LTV
- Monthly visitors, traffic sources
- Current marketing activities (SEO/social/ads/sales team)
- Known competitors (if user names them, skip discovery guesswork)
- USP, mission, vision
- Tech stack, CRM, analytics tools
- Brand guidelines / tone of voice

Data validation rules:

- If website URL provided → attempt automated enrichment (see §11, Integrations) before asking user to manually re-enter data already visible on the site.
 - If competitors not provided → system must research and propose 3-5 candidates for user confirmation before deep competitor analysis (prevents hallucinated competitor names).
-

8. Output: Report Structure

Final deliverable is a single structured report (in-app + exportable) with these sections, in order:

1. Executive Summary
2. Business Overview
3. Market Research & Industry Analysis
4. Customer Research (Personas)
5. Competitor Analysis
6. SWOT
7. PESTLE
8. Porter's Five Forces
9. Business Model Analysis
10. Positioning & Brand Strategy
11. Pricing Strategy
12. Go-To-Market Strategy
13. Marketing Strategy (channel-by-channel)
14. Content Strategy

- 15. Sales & Funnel Strategy
- 16. Customer Acquisition & Retention Strategy
- 17. AI Automation Recommendations
- 18. Budget Recommendation
- 19. 90-Day Execution Plan
- 20. 6-Month & 12-Month Roadmap
- 21. KPI Dashboard (measurable, with targets)
- 22. Risk Assessment
- 23. Final Recommendations (ranked by impact)

Formatting rules for every recommendation in the report:

- State the recommendation
- Explain **why** (reasoning/evidence)
- Explain **how** (implementation steps)
- Tag: Impact (High/Med/Low), Difficulty (High/Med/Low), Est. Cost, Timeframe (short-term <90 days vs long-term)
- Use tables for anything comparative (competitors, channels, roadmap)

9. Functional Requirements

#	Requirement	Priority
F1	Guided intake flow collecting schema in §7	P0
F2	Hard-gate logic: block strategy generation until required fields present	P0
F3	Web research module (live search) for market/industry/competitor data	P0
F4	Competitor auto-discovery when user doesn't name competitors, with user confirmation step	P1
F5	Full report generation following structure in §8	P0
F6	Export to PDF and Word (.docx)	P0
F7	Editable/regenerable sections (user can ask to redo "just the pricing strategy")	P1
F8	Save/resume engagements (multiple businesses per user account)	P1
F9	Source citation for every data-backed claim (industry size, trends, competitor facts)	P0

#	Requirement	Priority
F10	Assumption flagging — any claim not backed by user input or research must be labeled "Assumption"	P0
F11	KPI dashboard with measurable targets, not vague goals	P1
F12	Version history per engagement (strategy evolves over time)	P2
F13	Multi-business / agency mode (manage multiple client reports)	P2

10. Non-Functional Requirements

- **No hallucination policy:** Any statistic, competitor claim, or market figure must come from a retrieved source or be explicitly labeled as an estimate/assumption. This is a trust-critical requirement — violating it destroys product credibility.
- **Latency:** Full report generation is expected to take minutes, not seconds, given the research depth — set user expectations with progress indicators per step (§6).
- **Consistency:** Re-running the same inputs should produce strategically consistent (not contradictory) outputs across sections — e.g., pricing strategy must match the target audience's income bracket from customer research.
- **Auditability:** Every report should show which inputs were user-provided vs. AI-researched vs. AI-assumed.
- **Data privacy:** Business data (revenue, budget, customer data) is sensitive — needs standard encryption at rest/in transit and clear data-retention policy, especially for agency/multi-client use.

11. Technical Architecture Recommendations

High-level stack (dev to finalize based on team's existing infra):

- **LLM backbone:** A strong reasoning model (e.g., Claude) as the core "consultant brain," orchestrating the 10-step workflow via a structured prompt/state machine rather than one giant unstructured prompt — this keeps outputs consistent and gate-able.
- **Research/tooling layer:**
 - Live web search API (for market size, trends, news)
 - Website scraping/enrichment (for auto-filling business data from a URL)
 - Optional SEO data APIs (e.g., SEMrush/Ahrefs-style) for competitor SEO/backlink analysis — flag as a paid-tier feature given cost
 - Social listening APIs (optional, later phase) for competitor social performance

- **State management:** Each engagement is a persistent object with: intake data → research cache → generated report sections → version history. Recommend a document/JSON store per engagement rather than treating each chat turn as stateless.
 - **Report renderer:** Structured JSON output from the LLM → templated into the report UI + export formats (PDF/DOCX), rather than free-text parsing — far more reliable.
 - **Export:** DOCX/PDF generation pipeline (headings, tables, styled sections matching §8 structure).
-

12. UX/UI Requirements

- **Primary flow:** Guided form/chat hybrid intake → visible progress through the 10 steps (like a checklist, not a black box) → generated report in a readable document view.
 - **Section-level regeneration:** User can click "regenerate" on any single section without re-running the whole engagement.
 - **Trust indicators:** Small tags on claims — "Source: [research]" vs "Assumption — please confirm."
 - **Dashboard view:** For KPIs and roadmap — visual (table/timeline), not just paragraph text.
 - **Export button:** One click to PDF/DOCX.
 - **Multi-engagement management:** List view of past reports if user manages multiple businesses/clients.
-

13. MVP Scope (Phase 1) vs. Future Phases

Phase 1 (MVP):

- Guided intake (§7 required fields only)
- Steps 1–9 of workflow, using live web search only (no paid SEO API)
- Full report per §8, exportable as PDF
- Single-business per user, no multi-client mode

Phase 2:

- DOCX export, section regeneration
- Competitor auto-discovery with confirmation
- KPI dashboard visualization
- Save/resume + version history

Phase 3:

- Multi-client/agency mode
- Paid SEO/backlink data integration
- Website auto-enrichment from URL
- Automation recommendation execution (e.g., direct CRM/chatbot setup guidance with tool-specific steps)

14. Risks & Mitigations

Risk	Mitigation
LLM hallucinates market stats/competitor data	Mandatory citation + assumption-flagging (F9/F10); block unverified claims from final report
User skips intake and expects instant generic report	Hard-gate on required fields (F2)
Report feels "generic" despite structure	Force every recommendation to include why/how/impact/difficulty/cost tags (§8 formatting rules)
Research step latency frustrates users	Clear step-by-step progress UI, set expectations upfront
Over-scoped MVP delays launch	Follow phased scope in §13 strictly

15. Open Questions for Dev/Product Discussion

1. Chat-based intake, form-based intake, or hybrid — final decision affects backend state design significantly.
2. Which live search/data providers are budgeted for MVP vs. later phases?
3. Pricing model for the product itself (per-report, subscription, agency seats)?
4. Should section regeneration (F7) be in MVP or Phase 2 — affects prompt/state architecture complexity.
5. Data retention policy for sensitive client business data (especially for agency/multi-client mode).

End of document. This PRD should be treated as a living document — update as scope decisions are made with the dev team.